

HARSH VARDHAN SINGH

PRODUCT MANAGER · AI PRODUCTS

harshvsingh.work@gmail.com · linkedin.com/in/harshv5111 · github.com/Flukeshotz
portfolio-tau-ashen-94.vercel.app · Bengaluru · Open to relocation

SUMMARY

Product manager who finds high-friction workflows and ships AI systems that remove them — and can build the prototype himself. At HighRadius, self-initiated an SDR workflow tool and ran a leadership-sanctioned A/B test that delivered **27 net-new meetings and a 15% MCR lift** — the only intern to take a self-initiated MVP to a metric-validated experiment. Now in the Founder's Office at Skillcase owning app, GTM and content. Four AI products in production, each owned end-to-end: research → spec → build → ship → measure.

EXPERIENCE

Skillcase — Founder's Office · Product & App Owner

Jun 2026 – Present

AI-powered language-learning EdTech · Bengaluru (Hybrid)

- Own **app, GTM and content**. Replaced manual curriculum production with an end-to-end AI pipeline generating CEFR-aligned German lessons (Gemini + structured JSON), so content scales without adding headcount.
- Designed a **multi-stage QA engine** — deterministic validation, semantic similarity, vocabulary-overlap detection and a self-healing repair loop — that auto-regenerates failed lessons before human review, keeping a subject expert as final approver on everything learner-facing.
- Shipped three further generation pipelines to production (instructional imagery, exam-prep audio/video, structured exam papers), each gated on expert verification before release.
- Run continuous discovery — user interviews plus a standing feedback panel — converting findings into prioritised fixes, new-flow wireframes and content/social growth.

HighRadius — Product-Focused Intern

Sep 2025 – Jan 2026

Enterprise FinTech SaaS · Order-to-Cash · Fortune 500 clients · Hyderabad

- Diagnosed **4–5 tab context-switching — not skill gaps** — as the real SDR bottleneck through direct observation; built a single-screen workspace consolidating scripts, KPIs, ERP data and objection frameworks. Leadership-approved A/B test delivered **27 net-new meetings and a 15% MCR lift**. Sole intern to convert a self-initiated MVP into a sanctioned, metric-validated experiment; named POC for the intern cohort.
- Identified that **<15% of SDR calls were audited** — a coaching blind spot — and designed a 10-point LLM scoring system (temp = 0.0) with deterministic evaluation and human-in-the-loop routing, so the model never auto-decides on high-stakes calls. Prototyped a real-time WebSocket HUD for live pitch and speaking-rate feedback. Presented to Senior Manager, VP and CMO.

NextLeap PM Fellowship — PM Fellow

Apr – Jul 2026

Product research & strategy · Remote

- Reframed ChatGPT Voice's adoption problem from **transcription accuracy to interaction-model failure**, validated by 30+ surveys and 5 interviews — 65–70% of non-adopters cited loss of cognitive control as the #1 barrier. Built a KPI tree and prioritised UX roadmap targeting **2% → 6% voice share**, plus a V2V/V2T React prototype to validate the redesigned model before engineering investment.

AI PRODUCTS SHIPPED — built and deployed independently

PULSE — autonomous product-intelligence pipeline. Processes **8,000+ fintech app reviews/week** across four apps (Groww · INDmoney · Zerodha Kite · Stable Money) at **\$0 operating cost**; a 5-stage RapidFuzz validation layer proves every quote verbatim against source — zero hallucinated quotes in production. Built a cross-app **Compare view** that surfaces shared problem patterns and evidence-scored findings across all four. Surfaced a 0.75-star quality decline invisible to lifetime averages. · Python · UMAP · HDBSCAN · GitHub Actions · MCP

SIF Copilot — safety-first RAG platform indexing **2,500+ SEBI/AMFI document chunks**; cites every source and defers on high-stakes queries instead of answering with false certainty. · FastAPI · Qdrant · Groq · React

Gourmet AI — 3-agent pipeline (Ranker → Critic → Synthesizer) on Llama-3-70B; the Critic validates every output against a deterministic candidate list, while cost-aware routing and SHA-256 prompt caching cut token spend **~70%**. · Python · Groq

SELECTED PM CASE STUDIES

PawPal — consumer product strategy: JTBD, four-segment analysis, trust-first prioritisation, North Star metric, subscription model and defensibility.

Voice Input on ChatGPT India — market opportunity, UX teardown, competitive benchmark, KPI-tree roadmap.

SKILLS

Product — Problem framing · User research · PRD writing · A/B testing · RICE · Kano · KPI trees · Journey mapping · Prioritisation · Stakeholder management · Agile

AI / ML — LLMs (Gemini · Groq · Llama-3) · RAG · Prompt engineering · Multi-agent systems · Vector DBs (Qdrant · FAISS) · HDBSCAN / UMAP · Hallucination detection · Deterministic validation · Human-in-the-loop design · MCP

Technical — Python · SQL · FastAPI · React · Next.js · Power BI · Figma · Notion · GitHub · Vercel · AI-paired development (Claude · Antigravity) used to prototype, debug and ship every product above

EDUCATION

SRM Institute of Science & Technology — B.Tech CSE, AI/ML specialisation · CGPA 8.53 / 10

May 2026